

Predicting Sales Volume for 4 Different Product Types

Jennifer Brosnahan

7/10/2020

The purpose of this project is to predict sales of four different product types and assess the impact Service reviews and Customer reviews have on sales.

Predicted results on new dataset for four product types are shown at the end.

Target variable: 'Volume' for the product types: PC, Laptops, Netbooks, and Smartphones

Loading packages

```
library(tidyverse)
library(caret)
library(ggplot2)
library(corrplot)
library(openxlsx)
library(h2o)
library(kableExtra)
```

Importing data

```
existing <- read.csv(file.path('C:/Users/jlbro/OneDrive/C3T3', 'existing.csv'),
                     stringsAsFactors = TRUE)
```

Checking structure

```
str(existing)
```

```
## 'data.frame': 80 obs. of 18 variables:
## $ ProductType : Factor w/ 12 levels "Accessories",...: 7 7 7 5 5 1 1 1 1 1 ...
## $ ProductNum : int 101 102 103 104 105 106 107 108 109 110 ...
## $ Price : num 949 2250 399 410 1080 ...
## $ x5StarReviews : int 3 2 3 49 58 83 11 33 16 10 ...
```

```
## $ x4StarReviews      : int  3 1 0 19 31 30 3 19 9 1 ...
## $ x3StarReviews      : int  2 0 0 8 11 10 0 12 2 1 ...
## $ x2StarReviews      : int  0 0 0 3 7 9 0 5 0 0 ...
## $ x1StarReviews      : int  0 0 0 9 36 40 1 9 2 0 ...
## $ PositiveServiceReview: int  2 1 1 7 7 12 3 5 2 2 ...
## $ NegativeServiceReview: int  0 0 0 8 20 5 0 3 1 0 ...
## $ Recommendproduct    : num  0.9 0.9 0.9 0.8 0.7 0.3 0.9 0.7 0.8 0.9 ...
## $ BestSellersRank     : int 1967 4806 12076 109 268 64 NA 2 NA 18 ...
## $ ShippingWeight      : num 25.8 50 17.4 5.7 7 1.6 7.3 12 1.8 0.75 ...
## $ ProductDepth        : num 23.9 35 10.5 15 12.9 ...
## $ ProductWidth        : num 6.62 31.75 8.3 9.9 0.3 ...
## $ ProductHeight       : num 16.9 19 10.2 1.3 8.9 ...
## $ ProfitMargin        : num 0.15 0.25 0.08 0.08 0.09 0.05 0.05 0.05 0.05 0.05 ...
## $ Volume              : int 12 8 12 196 232 332 44 132 64 40 ...
```

Because regression algorithms can easily misinterpret categorical variables in which there are more than 2 values, we will dummify categorical data for regression modeling to binarize the values.

```
existingDummy <- dummyVars('~ .', data = existing)
existing2 <- data.frame(predict(existingDummy, newdata = existing))
```

Check structure again

```
str(existing2)
```

```
## 'data.frame': 80 obs. of 29 variables:
## $ ProductType.Accessories : num 0 0 0 0 0 1 1 1 1 1 ...
## $ ProductType.Display : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.ExtendedWarranty: num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.GameConsole : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Laptop : num 0 0 0 1 1 0 0 0 0 0 ...
## $ ProductType.Netbook : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.PC : num 1 1 1 0 0 0 0 0 0 0 ...
## $ ProductType.Printer : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.PrinterSupplies : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Smartphone : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Software : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Tablet : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductNum : num 101 102 103 104 105 106 107 108 109 110 ...
## $ Price : num 949 2250 399 410 1080 ...
## $ x5StarReviews : num 3 2 3 49 58 83 11 33 16 10 ...
## $ x4StarReviews : num 3 1 0 19 31 30 3 19 9 1 ...
## $ x3StarReviews : num 2 0 0 8 11 10 0 12 2 1 ...
## $ x2StarReviews : num 0 0 0 3 7 9 0 5 0 0 ...
## $ x1StarReviews : num 0 0 0 9 36 40 1 9 2 0 ...
## $ PositiveServiceReview : num 2 1 1 7 7 12 3 5 2 2 ...
## $ NegativeServiceReview : num 0 0 0 8 20 5 0 3 1 0 ...
## $ Recommendproduct : num 0.9 0.9 0.9 0.8 0.7 0.3 0.9 0.7 0.8 0.9 ...
## $ BestSellersRank : num 1967 4806 12076 109 268 ...
```

```
## $ ShippingWeight      : num  25.8 50 17.4 5.7 7 1.6 7.3 12 1.8 0.75 ...
## $ ProductDepth       : num  23.9 35 10.5 15 12.9 ...
## $ ProductWidth       : num  6.62 31.75 8.3 9.9 0.3 ...
## $ ProductHeight      : num  16.9 19 10.2 1.3 8.9 ...
## $ ProfitMargin       : num  0.15 0.25 0.08 0.08 0.09 0.05 0.05 0.05 0.05 0.05 ...
## $ Volume              : num  12 8 12 196 232 332 44 132 64 40 ...
```

Check summary for descriptive and NAs

```
summary(existing2)
```

```
## ProductType.Accessories ProductType.Display ProductType.ExtendedWarranty
## Min.      :0.000          Min.      :0.0000          Min.      :0.000
## 1st Qu.:0.000          1st Qu.:0.0000          1st Qu.:0.000
## Median :0.000          Median :0.0000          Median :0.000
## Mean    :0.325          Mean    :0.0625          Mean     :0.125
## 3rd Qu.:1.000          3rd Qu.:0.0000          3rd Qu.:0.000
## Max.    :1.000          Max.    :1.0000          Max.    :1.000
##
## ProductType.GameConsole ProductType.Laptop ProductType.Netbook ProductType.PC
## Min.      :0.000          Min.      :0.0000          Min.      :0.000          Min.      :0.00
## 1st Qu.:0.000          1st Qu.:0.0000          1st Qu.:0.000          1st Qu.:0.00
## Median :0.000          Median :0.0000          Median :0.000          Median :0.00
## Mean    :0.025          Mean    :0.0375          Mean     :0.025          Mean     :0.05
## 3rd Qu.:0.000          3rd Qu.:0.0000          3rd Qu.:0.000          3rd Qu.:0.00
## Max.    :1.000          Max.    :1.0000          Max.    :1.000          Max.    :1.00
##
## ProductType.Printer ProductType.PrinterSupplies ProductType.Smartphone
## Min.      :0.00          Min.      :0.0000          Min.      :0.00
## 1st Qu.:0.00          1st Qu.:0.0000          1st Qu.:0.00
## Median :0.00          Median :0.0000          Median :0.00
## Mean    :0.15          Mean     :0.0375          Mean     :0.05
## 3rd Qu.:0.00          3rd Qu.:0.0000          3rd Qu.:0.00
## Max.    :1.00          Max.    :1.0000          Max.    :1.00
##
## ProductType.Software ProductType.Tablet ProductNum Price
## Min.      :0.000          Min.      :0.0000          Min.      :101.0          Min.      : 3.60
## 1st Qu.:0.000          1st Qu.:0.0000          1st Qu.:120.8          1st Qu.: 52.66
## Median :0.000          Median :0.0000          Median :140.5          Median : 132.72
## Mean    :0.075          Mean     :0.0375          Mean     :142.6          Mean     : 247.25
## 3rd Qu.:0.000          3rd Qu.:0.0000          3rd Qu.:160.2          3rd Qu.: 352.49
## Max.    :1.000          Max.    :1.0000          Max.    :200.0          Max.    :2249.99
##
## x5StarReviews x4StarReviews x3StarReviews x2StarReviews
## Min.      : 0.0          Min.      : 0.00          Min.      : 0.00          Min.      : 0.00
## 1st Qu.: 10.0          1st Qu.: 2.75          1st Qu.: 2.00          1st Qu.: 1.00
## Median : 50.0          Median : 22.00          Median : 7.00          Median : 3.00
## Mean    : 176.2          Mean     : 40.20          Mean     : 14.79          Mean     : 13.79
## 3rd Qu.: 306.5          3rd Qu.: 33.00          3rd Qu.: 11.25          3rd Qu.: 7.00
## Max.    :2801.0          Max.    :431.00          Max.    :162.00          Max.    :370.00
##
## x1StarReviews PositiveServiceReview NegativeServiceReview Recommendproduct
```

```
## Min. : 0.00 Min. : 0.00 Min. : 0.000 Min. :0.100
## 1st Qu.: 2.00 1st Qu.: 2.00 1st Qu.: 1.000 1st Qu.:0.700
## Median : 8.50 Median : 5.50 Median : 3.000 Median :0.800
## Mean : 37.67 Mean : 51.75 Mean : 6.225 Mean :0.745
## 3rd Qu.: 15.25 3rd Qu.: 42.00 3rd Qu.: 6.250 3rd Qu.:0.900
## Max. :1654.00 Max. :536.00 Max. :112.000 Max. :1.000
##
## BestSellersRank ShippingWeight ProductDepth ProductWidth
## Min. : 1 Min. : 0.0100 Min. : 0.000 Min. : 0.000
## 1st Qu.: 7 1st Qu.: 0.5125 1st Qu.: 4.775 1st Qu.: 1.750
## Median : 27 Median : 2.1000 Median : 7.950 Median : 6.800
## Mean : 1126 Mean : 9.6681 Mean : 14.425 Mean : 7.819
## 3rd Qu.: 281 3rd Qu.:11.2050 3rd Qu.: 15.025 3rd Qu.:11.275
## Max. :17502 Max. :63.0000 Max. :300.000 Max. :31.750
## NA's :15
## ProductHeight ProfitMargin Volume
## Min. : 0.000 Min. :0.0500 Min. : 0
## 1st Qu.: 0.400 1st Qu.:0.0500 1st Qu.: 40
## Median : 3.950 Median :0.1200 Median : 200
## Mean : 6.259 Mean :0.1545 Mean : 705
## 3rd Qu.:10.300 3rd Qu.:0.2000 3rd Qu.: 1226
## Max. :25.800 Max. :0.4000 Max. :11204
##
```

Reveals 15 NA's for 'BestSellersRank'

Deleting BestSellersRank, only variable with NAs

```
existing2$BestSellersRank <- NULL
```

Correlation matrix of all variables


```
corrData <- cor(existing2)
```

Exporting correlation to excel

```
write.xlsx(corrData, file = "corrData.xlsx", row.names=TRUE)
write.xlsx(existing2, file = 'existing2.xlsx')
```

Viewing correlation heatmap, as you can see, it's unreadable with so many variables

```
corrplot(corrData)
```



Removing 5 Star since perfect correlation of 1 to target variable, risks overfitting. Also removing low correlated variables.

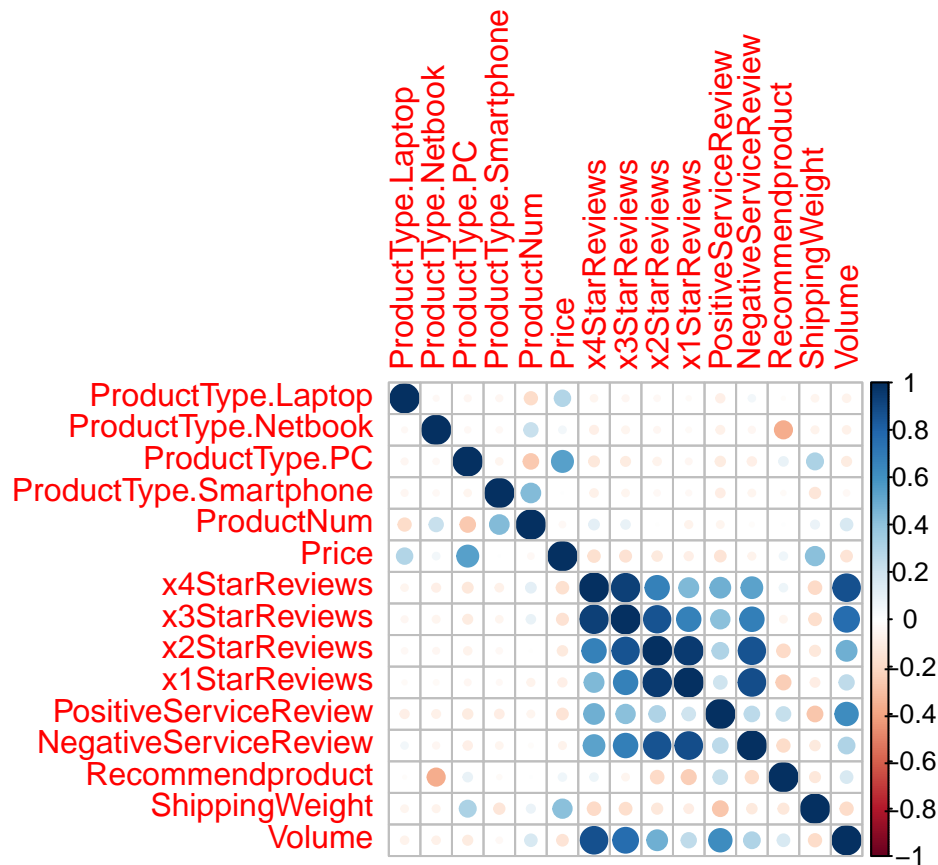
```
existing3 <- subset(existing2, select = -c(1:4, 8:9, 11:12, 15, 24:27))
str(existing3)
```

```
## 'data.frame': 80 obs. of 15 variables:
## $ ProductType.Laptop : num 0 0 0 1 1 0 0 0 0 0 ...
## $ ProductType.Netbook : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.PC : num 1 1 1 0 0 0 0 0 0 0 ...
## $ ProductType.Smartphone: num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductNum : num 101 102 103 104 105 106 107 108 109 110 ...
## $ Price : num 949 2250 399 410 1080 ...
## $ x4StarReviews : num 3 1 0 19 31 30 3 19 9 1 ...
## $ x3StarReviews : num 2 0 0 8 11 10 0 12 2 1 ...
## $ x2StarReviews : num 0 0 0 3 7 9 0 5 0 0 ...
## $ x1StarReviews : num 0 0 0 9 36 40 1 9 2 0 ...
## $ PositiveServiceReview : num 2 1 1 7 7 12 3 5 2 2 ...
## $ NegativeServiceReview : num 0 0 0 8 20 5 0 3 1 0 ...
## $ Recommendproduct : num 0.9 0.9 0.9 0.8 0.7 0.3 0.9 0.7 0.8 0.9 ...
## $ ShippingWeight : num 25.8 50 17.4 5.7 7 1.6 7.3 12 1.8 0.75 ...
## $ Volume : num 12 8 12 196 232 332 44 132 64 40 ...
```

EDA


viewing correlation heatmap

```
corrData3 <- cor(existing3)
corrplot(corrData3)
```



Enhancing the correlation heatmap. As you can see, x4Star, x3Star, x2Star, and PositiveService Review have highest correlation to target variable 'Volume.'


```
color <- colorRampPalette(c('#BB4444', '#EE9988', '#FFFFFF', '#77AADD', '#4477AA'))
corrplot(corrData3, method = 'shade', shade.col = NA, tl.col = 'black',
         type = 'upper', tl.srt = 45)
```



Histogram of Volume, reveals a couple outliers


```
ggplot(data = existing3, mapping = aes(x = Volume)) +  
  geom_histogram()
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



Plotting Sales Volume by Product Type. Our company is interested in sales volume for PCs, Laptops, Netbooks, and Smartphones for this project

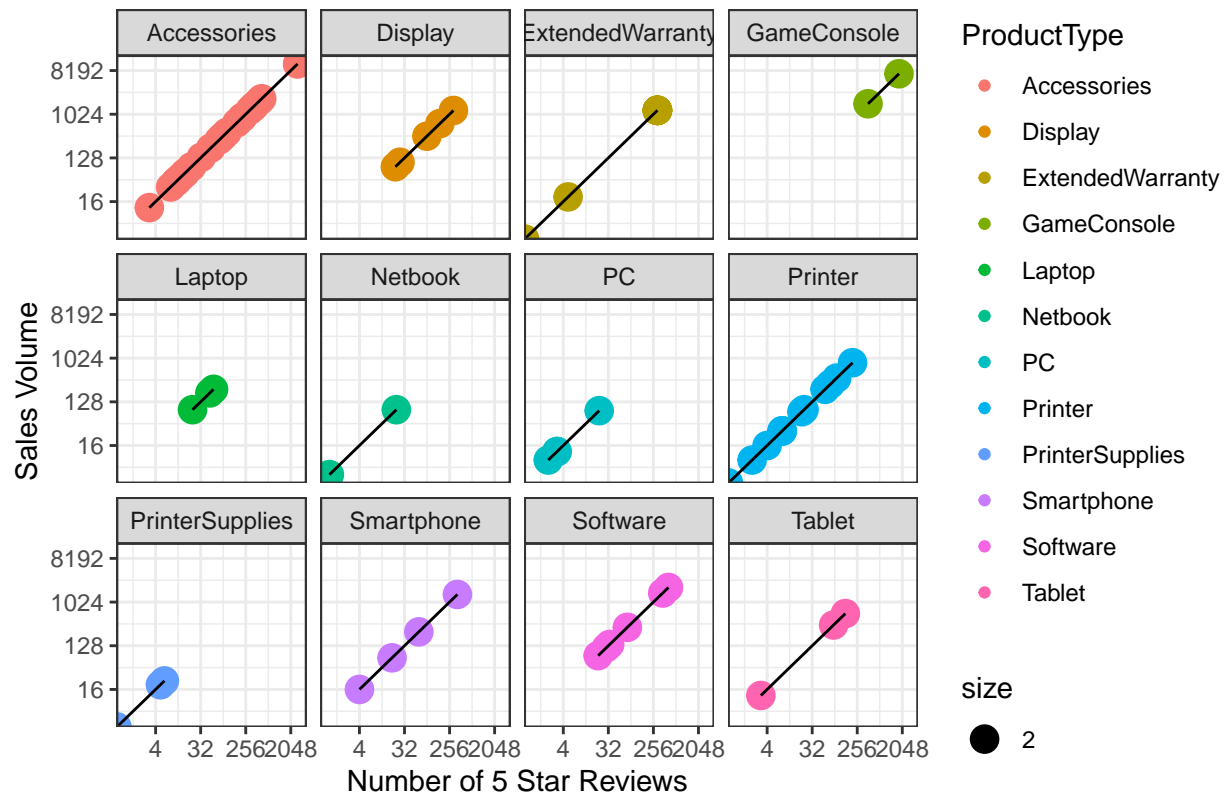
```
ggplot(data = existing, aes(x = ProductType, y = Volume, fill = ProductType)) +  
  geom_bar(stat = 'identity') +  
  guides(fill=FALSE) +  
  coord_flip()
```

Plotting the impact 5 Star Reviews have on Sales Volume. As you can see, it's a perfect correlation, which is impossible over time, thus why x5Star was removed from modeling

```
ggplot(data=existing, aes(x=x5StarReviews, y=Volume)) +
  geom_point(aes(color=ProductType, size=2)) +
  theme_bw() +
  scale_x_continuous(trans = 'log2') +
  scale_y_continuous(trans = 'log2') +
  geom_line() +
  facet_wrap(~ProductType) +
  xlab('Number of 5 Star Reviews') +
  ylab('Sales Volume') +
  ggtitle('Effect of 5 Star Reviews on Sales Volume')
```

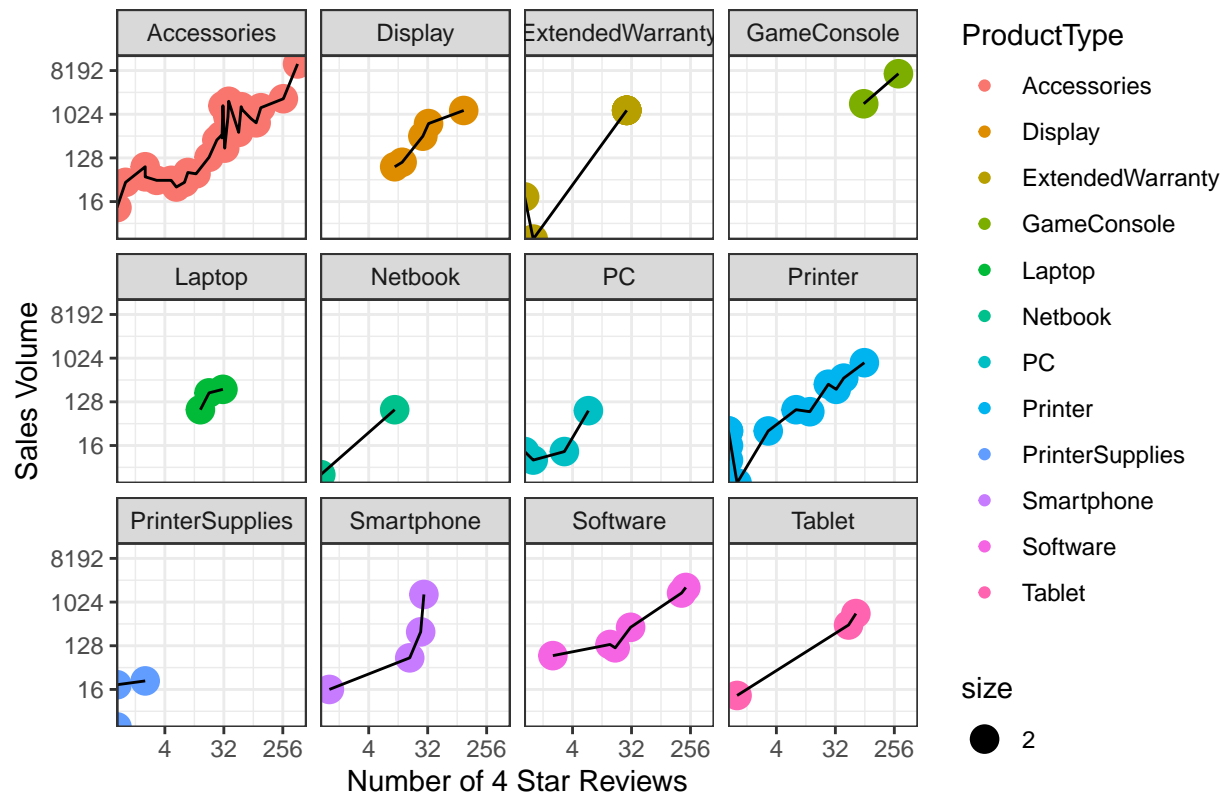
Effect of 5 Star Reviews on Sales Volume



Now plotting the impact of 4 Star Reviews on Sales Volume

```
ggplot(data=existing, aes(x=x4StarReviews, y=Volume)) +
  geom_point(aes(color=ProductType, size=2)) +
  theme_bw() +
  scale_x_continuous(trans = 'log2') +
  scale_y_continuous(trans = 'log2') +
  geom_line() +
  facet_wrap(~ProductType) +
  xlab('Number of 4 Star Reviews') +
  ylab('Sales Volume') +
  ggtitle('Effect of 4 Star Reviews on Sales Volume')
```


Effect of 4 Star Reviews on Sales Volume



Now plotting impact of Positive Service Reviews on Sales Volume

```
ggplot(data=existing, aes(x=PositiveServiceReview, y=Volume)) +
  geom_point(aes(color=ProductType, size=2)) +
  theme_bw() +
  scale_x_continuous(trans = 'log2') +
  scale_y_continuous(trans = 'log2') +
  geom_line() +
  facet_wrap(~ProductType) +
  xlab('Number of Positive Service Reviews') +
  ylab('Sales Volume') +
  ggtitle('Effect of Positive Service Reviews on Sales Volume')
```

Effect of Positive Service Reviews on Sales Volume



Modeling

Creating data partition and setting cross validation. Two rows eventually were removed due to outlier volumes.

```
set.seed(123)

# CreateDataPartition() 75% and 25%
index1 <- createDataPartition(existing3$Volume, p=0.75, list = FALSE)
train1 <- existing3[ index1,]
test1 <- existing3[-index1,]

# Removing 2 outlier rows #18 and #48 from test set
test1_rem_out <- test1[!rownames(test1) %in% c('18', '48'), ]

# Checking structure of train1
str(train1)

## 'data.frame': 61 obs. of 15 variables:
## $ ProductType.Laptop : num 0 0 0 1 0 0 0 0 0 0 ...
## $ ProductType.Netbook : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.PC : num 1 1 1 0 0 0 0 0 0 0 ...
## $ ProductType.Smartphone: num 0 0 0 0 0 0 0 0 0 0 ...
```

```
## $ ProductNum      : num  101 102 103 104 106 107 108 109 110 111 ...
## $ Price           : num  949 2250 399 410 114 ...
## $ x4StarReviews   : num   3  1  0 19 30  3 19  9  1  2 ...
## $ x3StarReviews   : num   2  0  0  8 10  0 12  2  1  2 ...
## $ x2StarReviews   : num   0  0  0  3  9  0  5  0  0  4 ...
## $ x1StarReviews   : num   0  0  0  9 40  1  9  2  0 15 ...
## $ PositiveServiceReview : num   2  1  1  7 12  3  5  2  2  2 ...
## $ NegativeServiceReview : num   0  0  0  8  5  0  3  1  0  1 ...
## $ Recommendproduct : num   0.9 0.9 0.9 0.8 0.3 0.9 0.7 0.8 0.9 0.5 ...
## $ ShippingWeight   : num  25.8 50 17.4 5.7 1.6 7.3 12 1.8 0.75 1 ...
## $ Volume           : num   12  8 12 196 332 44 132 64 40 84 ...
```

```
# Setting cross validation
control1 <- trainControl(method = 'repeatedcv',
                          number = 10,
                          repeats = 1)
```

Random forest model and tuning

```
# set seed
set.seed(123)

# Creating dataframe for manual tuning
rfGrid <- expand.grid(mtry = c(2,3,4,5,6,7,8))

rf1 <- train(Volume ~ x4StarReviews + PositiveServiceReview + x2StarReviews + x3StarReviews +
             x1StarReviews + NegativeServiceReview + Recommendproduct + ShippingWeight + Price,
             data = train1,
             method = 'rf',
             trControl = control1,
             tuneGrid = rfGrid)


rf1
```

```
## Random Forest
##
## 61 samples
## 9 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##  mtry  RMSE      Rsquared  MAE
##  2     869.2921  0.8755901  416.0646
##  3     849.5229  0.8871013  400.1741
##  4     824.7775  0.8939530  386.7741
##  5     827.7373  0.8980015  384.5929
##  6     801.6069  0.9043345  372.1979
##  7     802.9288  0.9073910  372.2425
```

```
##      8      788.6419  0.9081729  365.2084
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 8.
```

Level of importance for variables in model

```
ggplot(varImp(rf1, scale=FALSE)) +
  geom_bar(stat = 'identity', fill = 'steelblue') +
  ggtitle('Variable Importance of Random Forest 1 on Sales Volume')
```




Predicting rf on test1. Note, a symmetrical pattern means a good residual plot!

```
rf1Preds <- predict(rf1, newdata = test1_rem_out)
summary(rf1Preds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  14.11   19.05   102.97   425.15   719.71  1271.74
```

```
plot(rf1Preds)
```



Running a postResample to test if it will do well on new data or if overfitting. Our Cross Validation R2 is .908 after tuning and feature selection, which is excellent. Our postResample R2 is even better, at .945. If cross validation was above 94-95%, it would be a red-flag for overfitting, but postResample in upper 90s means it will generalize well on new data (and thus is not overfitting).

```
postResample(rf1Preds, test1_rem_out$Volume)
```

```
##          RMSE    Rsquared      MAE
## 190.4253816  0.9452884  98.7387608
```

CV RMSE=788, R2=.908

PostResample RMSE=190, R2=.945

Random Forest using feature selection

```
set.seed(123)

rf2 <- train(Volume ~ x4StarReviews + PositiveServiceReview + x2StarReviews,
             data = train1,
```

```
method = 'rf',  
trControl = control1)
```


note: only 2 unique complexity parameters in default grid. Truncating the grid to 2 .

```
rf2
```

```
## Random Forest  
##  
## 61 samples  
## 3 predictor  
##  
## No pre-processing  
## Resampling: Cross-Validated (10 fold, repeated 1 times)  
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...  
## Resampling results across tuning parameters:  
##  
##   mtry  RMSE      Rsquared  MAE  
##    2    771.2710  0.9218973  349.2301  
##    3    745.3771  0.9284383  338.8776  
##  
## RMSE was used to select the optimal model using the smallest value.  
## The final value used for the model was mtry = 3.
```

Variable importance


```
ggplot(varImp(rf2, scale=FALSE)) +  
  geom_bar(stat = 'identity', fill = 'steelblue') +  
  ggtitle('Variable Importance of Random Forest 2 on Sales Volume')
```

Plotting the residuals against the actual values for Volume. The graph below shows a couple volume outliers, and further research reveals both outliers are for accessories, which are not products of interest.

```
resid_rf2 <- residuals(rf2)
plot(train1$Volume, resid_rf2,
     xlab = 'Sales Volume',
     ylab = 'Residuals',
     main = 'Predicted Sales Volume Residuals Plot',
     abline(0,0))
```

Predicted Sales Volume Residuals Plot




Predicting rf2 on test1. This is another excellent residual plot, showing our predictions are consistent with regression.

```
rf2Preds <- predict(rf2, newdata = test1_rem_out)
summary(rf2Preds)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
##    2.742   14.336   103.066   419.753   714.428  1379.420
```

```
plot(rf2Preds)
```



postResample to test if it will do well on new data or if overfitting. It is even better than previous model (note: this is because we removed 2 outlier volumes from testSet, as indicated in DataPartition section above).

```
postResample(rf2Preds, test1_rem_out$Volume)
```

```
##          RMSE      Rsquared      MAE
## 153.8183544  0.9718405  74.7555212
```

CV RMSE = 745, R2=.928

PostResample RMSE=153, R2=.972

The postResample R2 and RMSE for a regression model is excellent. This is our top model!

Random Forest using feature selection

```
set.seed(123)

rf3 <- train(Volume ~ x4StarReviews + PositiveServiceReview + x3StarReviews,
             data = train1,
```

```
method = 'rf',  
trControl = control1)
```


note: only 2 unique complexity parameters in default grid. Truncating the grid to 2 .

```
rf3
```

```
## Random Forest  
##  
## 61 samples  
## 3 predictor  
##  
## No pre-processing  
## Resampling: Cross-Validated (10 fold, repeated 1 times)  
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...  
## Resampling results across tuning parameters:  
##  
##   mtry  RMSE      Rsquared  MAE  
##    2    710.5834  0.9285459  323.8009  
##    3    684.5642  0.9346194  313.5800  
##  
## RMSE was used to select the optimal model using the smallest value.  
## The final value used for the model was mtry = 3.
```

Variable importance

```
ggplot(varImp(rf3, scale=FALSE)) +  
  geom_bar(stat = 'identity', fill = 'steelblue') +  
  ggtitle('Variable Importance of Random Forest 3 on Sales Volume')
```



Predicting rf3 on test1

```
rf3Preds <- predict(rf3, newdata = test1_rem_out)
summary(rf3Preds)
```

```
##      Min.   1st Qu.   Median     Mean  3rd Qu.     Max.
##  5.789   11.959   94.221  425.461  737.133 1358.924
```

postResample to test if it will do well on new data or if overfitting. Another excellent model.

```
postResample(rf3Preds, test1_rem_out$Volume)
```

```
##      RMSE   Rsquared   MAE
## 167.7971733 0.9549701 83.4029171
```

CV RMSe=648, R2=.934

PostResample RMSE=167, R2=.954

Random Forest using feature selection


```
set.seed(123)
rf4 <- train(Volume ~ x4StarReviews + PositiveServiceReview + x3StarReviews + x2StarReviews +
             x1StarReviews + NegativeServiceReview,
             data = train1,
             method = 'rf',
             trControl = control1)
```

rf4

```
## Random Forest
##
## 61 samples
## 6 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##  mtry  RMSE      Rsquared  MAE
##  2      844.5227  0.8864850  396.6254
##  4      801.2553  0.9052961  370.6409
##  6      790.8997  0.9072861  362.8766
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 6.
```

Variable importance using ggplot

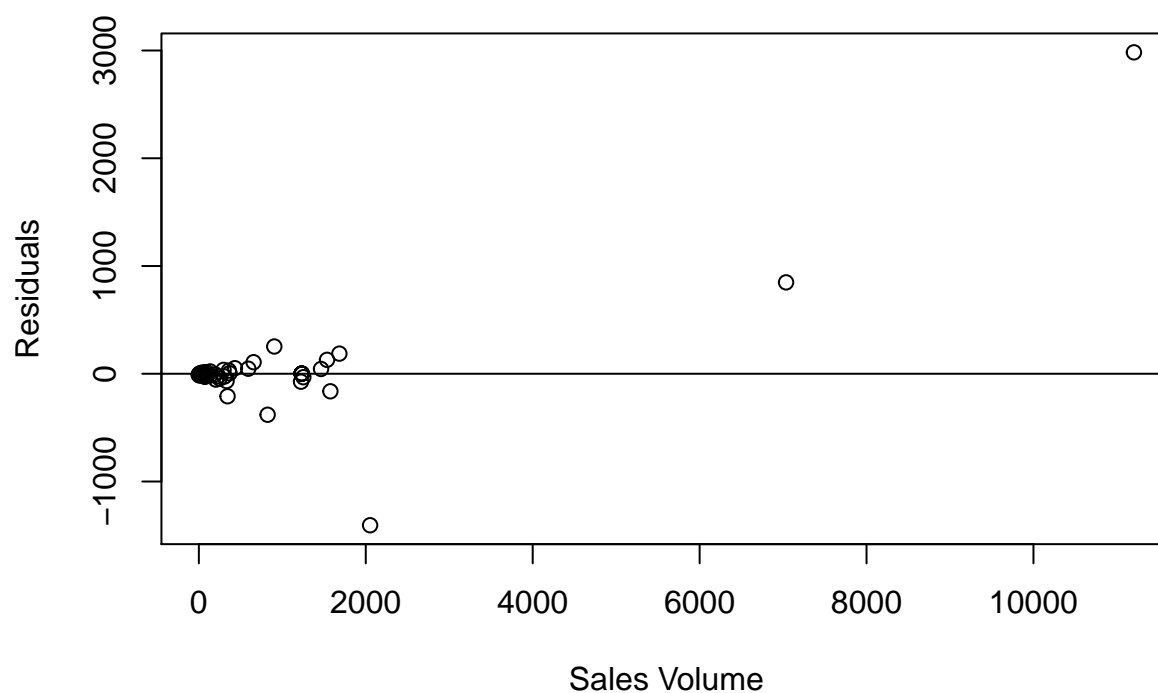
```
ggplot(varImp(rf4, scale=FALSE)) +
  geom_bar(stat = 'identity', fill = 'steelblue') +
  ggtitle('Variable Importance of Random Forest 4 on Sales Volume')
```



Plotting the residuals against the actual values for Volume. Again, graph shows outlier.

```
resid_rf4 <- residuals(rf4)
plot(train1$Volume, resid_rf4, xlab = 'Sales Volume', ylab = 'Residuals',
     main='Predicted Sales Volume Residuals Plot',
     abline(0,0))
```

Predicted Sales Volume Residuals Plot



Predicting rf4 on test1

```
rf4Preds <- predict(rf4, newdata = test1_rem_out)
summary(rf4Preds)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
##    7.937   15.732   98.618   426.150  782.476 1304.662
```

postResample to test if it will do well on new data or if overfitting

```
postResample(rf4Preds, test1_rem_out$Volume)
```

```
##      RMSE   Rsquared    MAE
## 177.453911  0.952956  87.506546
```


CV RMSE=783, R2=.909

RMSE=177, R2=.952

Support Vector Machines – RBF Kernel

```
set.seed(123)

# Creating dataframe for manual tuning
rbfGrid <- expand.grid(sigma = c(.01, .015, .2),
                      C = c(10, 100, 1000))

rbf1 <- train(Volume ~ x4StarReviews + x3StarReviews + PositiveServiceReview,
             data = train1,
             method = 'svmRadial',
             trControl = control1,
             tuneGrid = rbfGrid,
             preProc = c('center', 'scale'))

rbf1
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 61 samples
## 3 predictor
##
## Pre-processing: centered (3), scaled (3)
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##  sigma  C      RMSE      Rsquared  MAE
##  0.010  10    944.2228  0.8594778 489.1231
##  0.010  100   930.7863  0.8149199 473.1634
##  0.010  1000 1190.2695  0.8309208 580.6230
##  0.015  10    986.8673  0.8419784 507.9508
##  0.015  100   940.5278  0.8123009 480.0726
##  0.015  1000 1236.2804  0.8613697 590.2852
##  0.200  10    913.7802  0.9182651 467.5324
##  0.200  100   879.2184  0.9194906 436.6931
##  0.200  1000   949.7834  0.9093674 462.1128
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.2 and C = 100.
```

Predicting rbf on test1

```
rbf1Preds <- predict(rbf1, newdata = test1_rem_out)
summary(rbf1Preds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  29.81   94.91  267.83  463.59  552.12 2146.20
```

postResample to test if it will do well on new data or if overfitting

```
postResample(rbf1Preds, test1_rem_out$Volume)
```

```
##      RMSE      Rsquared      MAE
## 264.0730623  0.8148197 177.1172248
```

CV RMSE=879, R2=.919

PostResample RMSE=264, R2=.815

Support Vector Machines – RBF Kernel feature selection

```
set.seed(123)

# Creating dataframe for manual tuning
rbfGrid <- expand.grid(sigma = c(.01, .015, .2),
                      C = c(10, 100, 1000))

rbf2 <- train(Volume ~ x4StarReviews + PositiveServiceReview + x2StarReviews,
              data = train1,
              method = 'svmRadial',
              trControl = control1,
              tuneGrid = rbfGrid,
              preProc = c('center', 'scale'))

rbf2
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 61 samples
## 3 predictor
##
## Pre-processing: centered (3), scaled (3)
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##  sigma  C      RMSE      Rsquared  MAE
##  0.010  10   935.7611  0.8655808 465.7621
##  0.010  100  774.4498  0.8457538 408.1163
##  0.010 1000  729.0330  0.8563199 400.4417
##  0.015   10   879.8614  0.8620022 452.0967
##  0.015  100   761.7506  0.8470788 417.3052
##  0.015 1000   791.0228  0.8280795 418.6904
##  0.200   10   838.7191  0.9322844 436.6363
```

```
## 0.200 100 783.3714 0.9512888 398.3164
## 0.200 1000 657.2931 0.9099973 356.3912
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.2 and C = 1000.
```

Predicting rbf on test1

```
rbf2Preds <- predict(rbf2, newdata = test1_rem_out)
summary(rbf2Preds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -253.3    71.6    256.6    512.9    343.3    2499.2
```

```
# postResample to test if it will do well on new data or if overfitting
postResample(rbf2Preds, test1_rem_out$Volume)
```

```
##          RMSE      Rsquared      MAE
## 420.0433306  0.7040338 243.6208248
```

CV RMSE=657, R2=.909

PostResample RMSE=420, R2=.704

Negatives

Support Vector Machines – Linear

```
set.seed(123)

### Creating dataframe for manual tuning
linearGrid <- expand.grid(C = c(1, 10, 100, 1000))

linear1 <- train(Volume ~ x4StarReviews + x3StarReviews + PositiveServiceReview,
                 data = train1,
                 method = 'svmLinear',
                 trControl = control1,
                 tuneGrid = linearGrid,
                 preProc = c('center', 'scale'))

linear1
```

```
## Support Vector Machines with Linear Kernel
##
## 61 samples
## 3 predictor
##
## Pre-processing: centered (3), scaled (3)
```

```
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##      C      RMSE      Rsquared    MAE
##      1  873.4355  0.8664637  456.9757
##     10  843.0240  0.8585048  447.9310
##    100  848.6679  0.8571624  453.5503
##   1000  847.9066  0.8571915  452.4325
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was C = 10.
```

Predicting rbf on test1

```
linearPreds <- predict(linear1, newdata = test1_rem_out)
summary(linearPreds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -196.9  -129.0   124.7   388.2   561.6  2364.1
```

postResample to test if it will do well on new data or if overfitting

```
lin_PR <- postResample(linearPreds, test1_rem_out$Volume)
```

CV RMSE=843, R2=.858

PR RMSE=462, R2=.583

Negative predictions, move on

SVM – Linear, changing features

```
set.seed(123)

# Creating dataframe for manual tuning
linearGrid <- expand.grid(C = c(1, 10, 100, 1000))

linear2 <- train(Volume ~ x4StarReviews + x3StarReviews + PositiveServiceReview +
                  NegativeServiceReview + Price,
                  data = train1,
                  method = 'svmLinear',
                  trControl = control1,
                  tuneGrid = linearGrid,
                  preProc = c('center', 'scale'))

linear2
```

```
## Support Vector Machines with Linear Kernel
##
## 61 samples
## 5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##      C      RMSE      Rsquared    MAE
##      1  553.1088  0.8300446  328.9984
##     10  538.0562  0.8151764  327.8059
##    100  538.1027  0.8148848  327.7453
##   1000  538.5593  0.8146058  327.9203
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was C = 10.
```

```
# Predicting rbf on test1
linear2Preds <- predict(linear2, newdata = test1_rem_out)
summary(linear2Preds)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -155.3142 -102.7355   0.5217  435.3928  496.4102 2631.1579
```

postResample to test if it will do well on new data or if overfitting

```
postResample(linear2Preds, test1_rem_out$Volume)
```

```
##      RMSE      Rsquared      MAE
## 501.1605728  0.5858754 334.9077977
```

RMSE=1120, R2=56.9

Negative predictions, move on

Support Vector Machines – Polynomial

```
set.seed(123)

# Creating dataframe for manual tuning
polyGrid <- expand.grid(degree = c(2,3,4),
                        scale = c(1,2),
                        C = c(.1, 1, 10, 100))

poly1 <- train(Volume ~ x4StarReviews + x3StarReviews + PositiveServiceReview,
               data = train1,
```

```

method = 'svmPoly',
trControl = control1,
tuneGrid = polyGrid,
preProc = c('center','scale'))

```

```
poly1
```

```

## Support Vector Machines with Polynomial Kernel
##
## 61 samples
## 3 predictor
##
## Pre-processing: centered (3), scaled (3)
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##  degree  scale  C      RMSE      Rsquared  MAE
##  2        1      0.1    1155.408    0.7966076    571.6401
##  2        1      1.0    4104.003    0.8301827    1770.0867
##  2        1     10.0    6987.814    0.8557033    2939.2082
##  2        1    100.0    9796.456    0.8389599    4083.3127
##  2        2      0.1    2104.421    0.8267500     966.7124
##  2        2      1.0    6491.915    0.8763071    2740.5152
##  2        2     10.0    9245.414    0.8522459    3859.6825
##  2        2    100.0   10042.438    0.8331194    4185.3819
##  3        1      0.1    1118.828    0.8898300     534.3918
##  3        1      1.0    3755.882    0.9381422    1597.5451
##  3        1     10.0    7095.492    0.9176351    2965.7635
##  3        1    100.0   39627.409    0.9024923   16252.9238
##  3        2      0.1    2667.568    0.8509251    1166.7636
##  3        2      1.0    4472.345    0.8879559    1896.1439
##  3        2     10.0   31802.988    0.9179907   13053.9787
##  3        2    100.0   86413.899    0.8848359   35365.6943
##  4        1      0.1    6877.359    0.8003001    2881.4997
##  4        1      1.0   54000.769    0.8216935   22115.1070
##  4        1     10.0  197327.453    0.9111491   80637.6759
##  4        1    100.0  25145.120    0.9146607   10346.8229
##  4        2      0.1   31923.943    0.8869538   13115.3285
##  4        2      1.0  167176.011    0.8980690   68329.9630
##  4        2     10.0   68279.157    0.9005098   27957.4632
##  4        2    100.0  157877.892    0.9005079   64537.4960
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were degree = 3, scale = 1 and C = 0.1.

```

Predicting rbf on test1

```

polyPreds <- predict(poly1, newdata = test1_rem_out)
summary(polyPreds)

```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
## -31.78 -20.75 218.56 370.45 536.88 1240.00
```

postResample to test if it will do well on new data or if overfitting

```
postResample(polyPreds, test1_rem_out$Volume)
```

```
##          RMSE      Rsquared      MAE
## 334.4330248  0.7687032 174.1454650
```

RMSE=688, R2=60.2

Negative predictions, move on

SVM – Polynomial

```
set.seed(123)

# Creating dataframe for manual tuning
polyGrid <- expand.grid(degree = c(2,3,4),
                      scale = c(1,2),
                      C = c(.1, 1, 10, 100))

poly2 <- train(Volume ~ x4StarReviews + x2StarReviews + PositiveServiceReview +
               NegativeServiceReview,
               data = train1,
               method = 'svmPoly',
               trControl = control1,
               tuneGrid = polyGrid,
               preProc = c('center', 'scale'))

poly2
```

```
## Support Vector Machines with Polynomial Kernel
##
## 61 samples
## 4 predictor
##
## Pre-processing: centered (4), scaled (4)
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##  degree  scale  C      RMSE      Rsquared  MAE
##  2       1     0.1    876.0226  0.8081674  469.4007
##  2       1     1.0    1532.8601  0.8154286  722.1496
##  2       1    10.0   14512.8090  0.8528996  6032.6490
##  2       1   100.0   11848.1045  0.9129566  4942.3458
##  2       2     0.1    1115.2333  0.8751110  559.6023
##  2       2     1.0    5251.8265  0.8636670  2247.9786
```

```
##      2      2      10.0   11368.9766  0.8529120   4751.2822
##      2      2     100.0   11354.9053  0.9063491   4732.0598
##      3      1       0.1   18313.5053  0.8935899   7558.3279
##      3      1       1.0   24882.0057  0.8523866  10250.6328
##      3      1      10.0   61551.6914  0.9182948  25214.9606
##      3      1     100.0   36782.7066  0.8210000  15136.3743
##      3      2       0.1   37470.6589  0.8460612  15381.2988
##      3      2       1.0   34037.7172  0.8417435  13992.1891
##      3      2      10.0   62873.4712  0.8259141  25775.7687
##      3      2     100.0  100280.0776  0.8102818  41118.1724
##      4      1       0.1   96870.0477  0.8372587  39632.3275
##      4      1       1.0   9912.6724  0.8582925   4146.7389
##      4      1      10.0  174996.0882  0.8816298  71529.9157
##      4      1     100.0  334667.4580  0.7981803 136752.8839
##      4      2       0.1  233541.8543  0.8434558  95446.7978
##      4      2       1.0  404176.1594  0.8485852 165110.5747
##      4      2      10.0  704939.6023  0.8342341 287898.1043
##      4      2     100.0  790317.3179  0.8342340 322758.2980
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were degree = 2, scale = 1 and C = 0.1.
```

Predicting rbf on test1

```
poly2Preds <- predict(poly2, newdata = test1_rem_out)
summary(poly2Preds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    132.6   154.0   256.1   469.9   472.7  1710.2
```

postResample to test if it will do well on new data or if overfitting

```
postResample(poly2Preds, test1_rem_out$Volume)
```

```
##      RMSE    Rsquared    MAE
## 402.3116913  0.5699793 256.8209379
```

RMSE=402, R2=0.57

Gradient Boosting

```
set.seed(123)

gbm1 <- train(Volume ~ x4StarReviews + x2StarReviews + PositiveServiceReview,
              data = train1,
              method = 'gbm',
              trControl = control1,
              preProc = c('center', 'scale'))
```



```
gbm1
```

```
## Stochastic Gradient Boosting
##
## 61 samples
## 3 predictor
##
## Pre-processing: centered (3), scaled (3)
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##   interaction.depth  n.trees  RMSE      Rsquared  MAE
##   1                  50      1010.966  0.8249911  571.2535
##   1                  100      1054.100  0.8371555  585.4725
##   1                  150      1024.901  0.8667286  552.3807
##   2                   50      1010.350  0.8575585  557.7797
##   2                  100      1046.985  0.8593534  568.5074
##   2                  150      1053.486  0.8588568  578.0264
##   3                   50      1010.362  0.8472539  564.3698
##   3                  100      1038.869  0.8615773  564.2031
##   3                  150      1055.289  0.8588360  567.2579
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were n.trees = 50, interaction.depth =
## 2, shrinkage = 0.1 and n.minobsinnode = 10.
```

Predicting gbm on test1

```
gbmPreds <- predict(gbm1, newdata = test1_rem_out)
summary(gbmPreds)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
##    9.966   35.707   40.684   587.565 1340.268 2091.828
```

postResample to test if it will do well on new data or if overfitting

```
postResample(gbmPreds, test1_rem_out$Volume)
```

```
##      RMSE      Rsquared      MAE
## 266.4904990  0.9105057 172.6952417
```

Awesome step! provides comparison of predictions to actual within same DF!

```
compare_gbm1 <- data.frame(test1_rem_out,gbmPreds)
```

CV RMSE=1010, R2=.858

PostResample RMSE=266, R2=.911

Gradient Boosting

```
set.seed(123)

gbm2 <- train(Volume ~ x4StarReviews + x3StarReviews + PositiveServiceReview,
              data = train1,
              method = 'gbm',
              trControl = control1,
              preProc = c('center','scale'))
```

gbm2

```
## Stochastic Gradient Boosting
##
## 61 samples
## 3 predictor
##
## Pre-processing: centered (3), scaled (3)
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 55, 53, 55, 55, 55, 56, ...
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  RMSE      Rsquared  MAE
##  1                   50      1016.102  0.8326938  568.3590
##  1                   100      1063.431  0.8452965  576.8135
##  1                   150      1039.162  0.8622175  553.1044
##  2                   50      1021.436  0.8660802  555.2569
##  2                   100      1041.471  0.8672959  564.3505
##  2                   150      1053.235  0.8573396  580.2639
##  3                   50      1009.603  0.8574569  552.9631
##  3                   100      1056.474  0.8524858  573.3080
##  3                   150      1066.260  0.8481968  572.4470
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were n.trees = 50, interaction.depth =
## 3, shrinkage = 0.1 and n.minobsinnode = 10.
```

Predicting gbm2 on test1

```
gbm2Preds <- predict(gbm2, newdata = test1_rem_out)
summary(gbm2Preds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -13.48   48.78   48.78  541.48 1105.80 2141.98
```

postResample to test if it will do well on new data or if overfitting

```
postResample(gbm2Preds, test1_rem_out$Volume)
```

```
##      RMSE    Rsquared    MAE
## 264.0765350  0.8822213 156.1233574
```

CV RMSE=813, R2=.962

PostResample RMSE=415, R2=.706

Bayesian Ridge Regression, L1

```
set.seed(123)

bay1 <- train(Volume ~ x4StarReviews + PositiveServiceReview + x2StarReviews,
              data = train1,
              method = 'blassoAveraged',
              trControl = control1,
              preProc = c('center', 'scale'))

bay1
```

Predicting bay1 on test1

```
bay1Preds <- predict(bay1, newdata = test1_rem_out)
summary(bay1Preds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -211.99 -187.44   79.51  397.59  667.00 2491.79
```

postResample to test if it will do well on new data or if overfitting

```
postResample(bay1Preds, test1$Volume)
```

```
##      RMSE Rsquared    MAE
## 925.6151      NA 716.4564
```

Negative predictions regardless of feature selection, high RMSE, doesn't work with this task

CV RMSE=1148, R2=.753

After deleting problem outlier rows in test set - 17 observations

```
Actual_vs_Predicted_NoOutlier <- data.frame(test1_rem_out %>% select(ProductNum, Volume),
rf1Preds, rf2Preds, rf3Preds, rf4Preds, rbf1Preds, rbf2Preds, linearPreds, linear2Preds, polyPreds,
poly2Preds, gbmPreds, gbm2Preds)
```

exporting to excel

```
write.xlsx(Actual_vs_Predicted_NoOutlier, file = "Actual_vs_Predicted_NoOutlier.xlsx", row.names=TRUE)
```

Using Top Model rf2 algorithm to make predictions on new product data

Target variable: 'Volume' for PC, Laptops, Netbooks, and Smartphones product types

Importing data

```
new <- read.csv(file.path('C:/Users/jlbro/OneDrive/C3T3', 'new.csv'), stringsAsFactors = TRUE)
```

Checking structure

```
str(new)
```

```
## 'data.frame': 24 obs. of 18 variables:
## $ ProductType : Factor w/ 12 levels "Accessories",...: 7 7 5 5 5 6 6 6 6 12 ...
## $ ProductNum : int 171 172 173 175 176 178 180 181 183 186 ...
## $ Price : num 699 860 1199 1199 1999 ...
## $ x5StarReviews : int 96 51 74 7 1 19 312 23 3 296 ...
## $ x4StarReviews : int 26 11 10 2 1 8 112 18 4 66 ...
## $ x3StarReviews : int 14 10 3 1 1 4 28 7 0 30 ...
## $ x2StarReviews : int 14 10 3 1 3 1 31 22 1 21 ...
## $ x1StarReviews : int 25 21 11 1 0 10 47 18 0 36 ...
## $ PositiveServiceReview: int 12 7 11 2 0 2 28 5 1 28 ...
## $ NegativeServiceReview: int 3 5 5 1 1 4 16 16 0 9 ...
## $ Recommendproduct : num 0.7 0.6 0.8 0.6 0.3 0.6 0.7 0.4 0.7 0.8 ...
## $ BestSellersRank : int 2498 490 111 4446 2820 4140 2699 1704 5128 34 ...
## $ ShippingWeight : num 19.9 27 6.6 13 11.6 5.8 4.6 4.8 4.3 3 ...
## $ ProductDepth : num 20.63 21.89 8.94 16.3 16.81 ...
## $ ProductWidth : num 19.2 27 12.8 10.8 10.9 ...
## $ ProductHeight : num 8.39 9.13 0.68 1.4 0.88 1.2 0.95 1.5 0.97 0.37 ...
## $ ProfitMargin : num 0.25 0.2 0.1 0.15 0.23 0.08 0.09 0.11 0.09 0.1 ...
## $ Volume : int 0 0 0 0 0 0 0 0 0 0 ...
```

Making new dataframe same column wise as trained dataframes

```
newDummy <- dummyVars('~ .', data = new)
```

```
new2 <- data.frame(predict(newDummy, newdata = new))
```

check structure again

```
str(new2)
```

```
## 'data.frame': 24 obs. of 29 variables:
## $ ProductType.Accessories : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Display : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.ExtendedWarranty: num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.GameConsole : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Laptop : num 0 0 1 1 1 0 0 0 0 0 ...
## $ ProductType.Netbook : num 0 0 0 0 0 1 1 1 1 0 ...
## $ ProductType.PC : num 1 1 0 0 0 0 0 0 0 0 ...
## $ ProductType.Printer : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.PrinterSupplies : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Smartphone : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Software : num 0 0 0 0 0 0 0 0 0 0 ...
## $ ProductType.Tablet : num 0 0 0 0 0 0 0 0 0 1 ...
## $ ProductNum : num 171 172 173 175 176 178 180 181 183 186 ...
## $ Price : num 699 860 1199 1199 1999 ...
## $ x5StarReviews : num 96 51 74 7 1 19 312 23 3 296 ...
## $ x4StarReviews : num 26 11 10 2 1 8 112 18 4 66 ...
## $ x3StarReviews : num 14 10 3 1 1 4 28 7 0 30 ...
## $ x2StarReviews : num 14 10 3 1 3 1 31 22 1 21 ...
## $ x1StarReviews : num 25 21 11 1 0 10 47 18 0 36 ...
## $ PositiveServiceReview : num 12 7 11 2 0 2 28 5 1 28 ...
## $ NegativeServiceReview : num 3 5 5 1 1 4 16 16 0 9 ...
## $ Recommendproduct : num 0.7 0.6 0.8 0.6 0.3 0.6 0.7 0.4 0.7 0.8 ...
## $ BestSellersRank : num 2498 490 111 4446 2820 ...
## $ ShippingWeight : num 19.9 27 6.6 13 11.6 5.8 4.6 4.8 4.3 3 ...
## $ ProductDepth : num 20.63 21.89 8.94 16.3 16.81 ...
## $ ProductWidth : num 19.2 27 12.8 10.8 10.9 ...
## $ ProductHeight : num 8.39 9.13 0.68 1.4 0.88 1.2 0.95 1.5 0.97 0.37 ...
## $ ProfitMargin : num 0.25 0.2 0.1 0.15 0.23 0.08 0.09 0.11 0.09 0.1 ...
## $ Volume : num 0 0 0 0 0 0 0 0 0 0 ...
```

```
new2$BestSellersRank <- NULL
```

```
str(new2)
```

```
new3 <- subset(new2, select = -c(1:4, 8:9, 11:12, 15, 24:27))
```

```
str(new3)
```

Predicting rbf1 on ‘new3’ product data

```
set.seed(123)

Predicted_Volume <- predict(rf2, newdata = new3)
```

Adding our predictions to the ‘new’ product dataframe

```
Preds_rf2_df <- data.frame(new3 %>% select(ProductType.Laptop, ProductType.Netbook, ProductType.PC, ProductType.Smartphone))
TopModelPreds <- read.xlsx(file.path('C:/Users/jlbro/OneDrive/C3T3-3', 'newPreds_TopModel_rf2.xlsx'))
```

Finally viewing our Sales predictions for 4 product types on a new dataset! Also showing ‘x4StarReviews’ since it was the most important variable for the model. All products not of interest have been removed from table.

```
kable(TopModelPreds) %>%
  kable_styling(bootstrap_options = c('striped', 'hover'))
```

Product.Type	ProductNum	4.Star.Reviews	Predicted_Volume
PC	171	26	478.64227
PC	172	11	157.28747
Laptop	173	10	187.16573
Laptop	175	2	36.68747
Laptop	176	1	14.43680
Netbook	178	8	55.57160
Netbook	180	112	1234.30893
Netbook	181	18	129.49760
Netbook	183	4	19.38773
Smartphone	193	26	444.73333
Smartphone	194	26	649.95707
Smartphone	195	8	87.20040
Smartphone	196	19	159.08307

Exporting to excel

```
write.xlsx(Preds_rf2_df, "TopModel_rf2_Preds.xlsx")
```